



Automated Artwork **Metadata Generation** using Multimodal LLM

Andò S. - Baldi F. - Cozzone R.

CHALLENGES

CONTEXT



Cultural institutions are experiencing a significant expansion in the digitization of their archives. Museums and art galleries, in particular, are working to convert extensive collections of artworks into digital formats that can be more easily preserved, searched, and accessed.

STAKEHOLDERS



Museum
Curators

Manual annotation is:

- time-consuming;
- expensive;
- non-scalable.



AI Researchers

- Limited availability of high-quality labeled datasets;
- Lack of standardized metadata schemas across institutions.

CHALLENGES

Reduce the time and effort required for manual metadata compilation

Overcome the scarcity of reliable labeled data and the inconsistency of metadata structures across different cultural institutions

OBJECTIVES



Allowing curators to focus on **higher-value activities** through the automatic generation of metadata.



Enriching a database with more detailed information **improves searchability**, facilitates research, and enhances the understanding of the content.




Standardize data structures to enable easier **information sharing** and improve public access to cultural heritage.



HYPOTHESES

Q: Can a multimodal LLM correctly provide informations about a painting?

A: MLLM can effectively bridge the gap between raw visual data and structured knowledge



DATASET




artefact	type	title	description	genres	materials	subjects	authors	images
----------	------	-------	-------------	--------	-----------	----------	---------	--------

https://wikidata	Painting	Madonna del Prato	painting by Raphael	['religius art']	['oil paint', 'poplar panel']	['woman', 'boy', 'Mary', 'Christ Child',..]	['Raphael']	02992163-d6..
---	----------	-------------------	---------------------	------------------	-------------------------------	---	-------------	---------------



Column	Unique elements		Total occurrences		Average elements/row		Max elements/row	
	TRAIN	VAL	TRAIN	VAL	TRAIN	VAL	TRAIN	VAL
Authors	3.282	1.040	7.372	1.773	1,02	1,01	9	2
Genres	87	41	7.601	1.882	1,05	1,07	5	4
Materials	110	47	14.047	3.494	1,94	1,99	6	5
Subjects	6.217	2.442	20.497	6.446	2,83	3,67	69	76



DATASET

PREDOMINANT CLASSES


There is a clear dominance of a few labels compared to the entire vocabulary.

- In material, terms like "**oil paint**" or "**canvas**" cover a significant portion of the dataset.
- In genres, categories like "**portrait**" or "**religious art**" dominate over niche genres.


IMPACT ON MODELING

This imbalance introduces a risk of **bias**: the model may learn the dominant classes perfectly but fail to generalize on the rare ones (overfitting on dominant classes).

GENRES



MATERIALS



METHODOLOGY



ZERO-SHOT EVALUATION

Evaluating the models "**As-Is**" (Pre-trained weights) to benchmark their intrinsic capabilities without initial fine-tuning.

- **VLM SELECTION**

We used **BLIP-2** and **Qwen2-VL** as our core Vision-Language Models to test varying architectures (Q-Former vs. Native Resolution).




FINE-TUNING

Fine-tune the model that demonstrates the highest baseline performance in benchmarks

- **PEFT TECHNIQUES**

Updates only a small subset of parameters avoiding to retrain the entire model.

QWEN2-VL-2B-INS.




CHAT-TEMPLATE

NATIVE RESOLUTION

**MULTIMODAL ROTARY
POSITION EMBEDDING**

BLIP-2 FLAN-T5-XL



INSTRUCTION-TUNED

FROZEN

**INFORMATION
BOTTLENECK**

METHODOLOGY



ZERO-SHOT EVALUATION

Evaluating the models "**As-Is**" (Pre-trained weights) to benchmark their intrinsic capabilities without initial fine-tuning.

- **VLM SELECTION**

We used **BLIP-2** and **Qwen2-VL** as our core Vision-Language Models to test varying architectures (Q-Former vs. Native Resolution).



FINE-TUNING

Fine-tune the model that demonstrates the highest baseline performance in benchmarks

- **PEFT TECHNIQUES**

Updates only a small subset of parameters avoiding to retrain the entire model.


METHODOLOGY

PEFT (Parameter-Efficient Fine-Tuning)

Family of techniques that adapt large pre-trained models to new tasks by training only a small subset of parameters.

- **BENEFITS**

Reduces memory usage, training time, and computational cost while preserving most of the model's performance.



LoRA (Low-Rank Adaptation)


PEFT method that injects low-rank matrices into selected layers (typically attention layers).

- **DETAILS**

Only low-rank parameters are trained, while the original model weights remain frozen.

- **OUR CONFIGURATIONS**

- Rank = 8 - Dropout = 0.1
- Rank = 16 - Dropout = 0.05



METHODOLOGY

LoRA TARGET MODULES

- **Vision-Language Adapter (`visual.merger.mlp`)**
Updates the Multi-Layer Perceptron within the visual merger.
- **Self-Attention Layers (`self.attn.`)**
Updates the Query, Key, Value and Output projections within the attention mechanism.
- **Feed-Forward Network (`mlp.`)**
 - Gate → filters the information
 - Up → maps inputs into a higher-dimensional space to extract complex features.
 - Down → compresses processed data back to the original space to merge with the network

```
lora_config = LoraConfig(  
    r=16, lora_alpha=32,  
    target_modules=[  
        "visual.merger.mlp.0",  
        "visual.merger.mlp.2",  
        "model.layers.*.self_attn.q_proj",  
        "model.layers.*.self_attn.k_proj",  
        "model.layers.*.self_attn.v_proj",  
        "model.layers.*.self_attn.o_proj",  
        "model.layers.*.mlp.gate_proj",  
        "model.layers.*.mlp.up_proj",  
        "model.layers.*.mlp.down_proj"],  
    lora_dropout=0.05, bias="none",  
    task_type=TaskType.CAUSAL_LM)
```

LOSS FUNCTION

- **Cross Entropy (`CAUSAL_LM`)**
 - Standard Next-Token Prediction loss used to train LMs.
 - Measures the discrepancy between prediction and GT.
 - Penalizes incorrect token predictions in a sequence

```
pred: [ "Subject:", "man" ]  
gt:   [ "Subject:", "woman" ]  
loss: [ low      , high  ]
```

METHODOLOGY

SINGLE-MODEL

1 SINGLE PROMPT

A holistic approach where the model is asked to predict all **four categories** within a **single inference cycle**.

Prompt:

"Identify the author(s), genre(s), material(s), and subject(s) of this artwork. Answer following this exact format: 'author: name | genre: name | material: name | subject: name'."



Multimodal
AI

OUTPUT

Authors : Parmigianino

Subjects: man, book

Materials: oil paint, panel

Genres: portrait

METHODOLOGY

SINGLE-MODEL

2 1 IMG - 4 PROMPT

Four independent **prompts** are generated for each image, with each prompt focusing on a **single** metadata **category**

Prompt **authors**:

Identify the author(s) of this artwork.
Answer with the format 'author: name'.



Prompt **subjects**:

Identify the subject(s) of this artwork.
Answer with the format 'subject: name'.



Prompt **materials**:

"Identify the material(s) used in this artwork.
Answer with the format 'material: name'."



Prompt **Genres**:

Identify the genre(s) of this artwork.
Answer with the format 'genre: name'.



Multimodal
AI

OUTPUT

Authors : Parmigianino

Subjects: man, book

Materials: oil paint, panel

Genres: portrait

METHODOLOGY

MULTI-MODEL

3 SINGLE PROMPT

Four distinct versions of the **model** are fine-tuned independently, each dedicated to a **single** metadata **category**

OUTPUT

Prompt **authors**:

Identify the author(s) of this artwork.
Answer with the format 'author: name'.



Multimodal
AI



Authors : Parmigianino

Prompt **subjects**:

Identify the subject(s) of this artwork.
Answer with the format 'subject: name'.



Multimodal
AI



Subjects: man, book

Prompt **materials**:

"Identify the material(s) used in this
artwork. Answer with the format
'material: name'."



Multimodal
AI



Materials: oil paint, panel

Prompt **Genres**:

Identify the genre(s) of this artwork.
Answer with the format 'genre: name'.



Multimodal
AI



Genres: portrait

EXPERIMENTAL SETUP

HYBRID TASK

- **QA problem:** VLM generate an answer as close as possible to the expected Wikidata answer.
- **Open-domain vocabulary:** the possible classes are not defined in advance.


SYNTACTIC EVALUATION

How many terms are correctly predicted? We cannot use standard precision, recall, F1 → **EXACT MATCH**

SEMANTIC EVALUATION

Exact Match cannot capture semantic similarity → **MPNET EMBEDDINGS**

⚠ Similarity('person', any author) ≥ 0.7



QA METRICS

- **BLEU** and **ROUGE** useful for the N-gram overlap analysis.
- These metrics expect answers with more than one token. Particularly, BLEU penalizes answer with only one token.



EXACT MATCH EXTENSION

Querying Wikidata database to retrieve aliases of authors



WIKIDATA




RESULTS

ZERO-SHOT



	Exact Match (Accuracy)		Similarity Score		BLEU		ROUGE-1	
	QWEN	BLIP	QWEN	BLIP	QWEN	BLIP	QWEN	BLIP
Authors	0.150	0.001	0.800	0.684	X	X	0.154	0.003
Subjects	0.014	0.000	0.717	0.722	0.030	0.014	0.193	0.099
Materials	0.218	0.000	0.886	0.838	0.210	0.105	0.609	0.478
Genres	0.064	0.353	0.586	0.549	0.079	0.077	0.324	0.485

QWEN WIN



RESULTS SINGLE-MODEL

1 IMG-4PROMPT

Exact Match (Accuracy)

	Authors	Genres	Materials	Subjects
R=8, D=0.1	17,71%	75,80%	68,74%	7,92%
R=16, D=0.05	18,45%	76,37%	69,19%	7,74%

Exact Match (Accuracy)

	Scheduler	Authors	Genres	Materials	Subjects
R=16, D=0.05	✗	18,45%	76,37%	69,19%	7,74%
	✓	18,28%	75,68%	69,59%	8,14%



R=16, D=0.05 & SCHEDULER



SINGLE PROMPT

Exact Match (Accuracy)

	Authors	Genres	Materials	Subjects
R=8, D=0.1	17,26%	75,34%	66,57%	6,66%
R=16, D=0.05	17,06%	75,57%	66,86%	7,00%

Exact Match (Accuracy)

	Scheduler	Authors	Genres	Materials	Subjects
R=16, D=0.05	✗	17,06%	75,57%	66,86%	7,00%
	✓	17,43%	75,97%	67,43%	6,66%



R=16, D=0.05 & SCHEDULER

RESULTS SINGLE-MODEL



1 IMG

PROMPT

	Authors
R=8, D=0.1	17,43%
R=16, D=0.05	18,28%

	Scheduler
R=16, D=0.05	✗
	✓

	Exact Match (Accuracy)				
	Strategy	Authors	Genres	Materials	Subjects
R=16, D=0.05 Scheduler	1 IMG-4PROMPT	18,28%	75,68%	69,59%	8,14%
	PROMPT UNICO	17,43%	75,97%	67,43%	6,66%

1 IMG-4PROMPT (R=16, D=0.05 & SCHEDULER)



	Materials	Subjects
66,57%	66,57%	6,66%
66,86%	66,86%	7,00%

	Materials	Subjects
66,57%	66,86%	7,00%
67,97%	67,43%	6,66%



R=16, D=0.05 & SCHEDULER



R=16, D=0.05 & SCHEDULER

RESULTS MULTIMODEL

UPPER BOUNDS

The Multi-Model performance are used as our **upper bound**.

BENCHMARKING

Benchmark our **Best Configuration** (with Single-Model strategy) against this to validate its competitiveness and efficiency.

Exact Match (Accuracy)

	Authors	Genres	Materials	Subjects
MULTI-MODEL	18,45%	76,65%	70,44%	7,00%
BEST CONFIGURATION	18,28%	75,68%	69,59%	8,14%

-0,17%

-0,97%

-0,85%

+1,14%

RESULTS

FINAL COMPARISON

BEST CONFIGURATION vs. ZERO-SHOT

A comparison between the best fine-tuned configuration and the zero-shot performance baseline **to quantify the performance gains**, considering all metrics.

Exact Match (Accuracy)

Similarity

	Authors	Genres	Materials	Subjects	Authors	Genres	Materials	Subjects
ZERO-SHOT BASELINE	15,00%	6,40%	21,80%	1,40%	80,00%	58,60%	88,60%	71,70%
BEST FINE-TUNED MODEL	18,28%	75,68%	69,59%	8,14%	86,31%	75,68%	96,43%	60,34%
	+3,28%	+69,28%	+47,79%	+6,26%	+6,31%	+17,08%	+7,83%	-11,36%

BLEU

ROUGE-1

	Authors	Genres	Materials	Subjects	Authors	Genres	Materials	Subjects
ZERO-SHOT BASELINE	x	7,90%	21,00%	3,00%	15,40%	32,40%	60,90%	19,30%
BEST FINE-TUNED MODEL	x	20,07%	44,60%	4,63%	21,25%	80,56%	87,33%	24,99%
	-	+12,17%	+23,60%	+1,63%	+6,15%	+48,16%	+26,43%	+5,69%

Mean percentage increase

Exact Match +31,65%

Similarity +4,96%

BLEU +12,47%

ROUGE-1 +21,61%


RESULTS

IMPACT OF DATASET IMBALANCE ON PREDICTIONS


Despite the imbalance in our training data, the model demonstrates low bias and avoided overfitting. It accurately reflects the actual distribution of labels without excessively predicting the dominant classes.

Ground Truth
Prediction

GENRES



MATERIALS



QUALITATIVE ASSESSMENT



	Authors	Materials	Genres	Subjects
Ground True	Jacob van Walscapelle	oil paint, panel	still life	pomegranate, fruit, grape, wine glass
Prediction	Jan Saenredam	oil paint, panel	still life	pomegranate, grape
				



	Authors	Materials	Genres	Subjects
Ground True	Giovanni Bellini	oil paint, panel	religious art	man, Jesus Christ
Prediction	Giovanni Bellini	oil paint, panel	religious art	Christ, Christ Pantocrator
				

CONCLUSIONS



- Approaches, like Parameter-Efficient Fine-Tuning (PEFT) and Low-Rank Adaptation (LoRA), improved model performance compared to the zero-shot baseline.
- Good results for “genres” and “materials” columns, achieving an exact match accuracy of 69,28% and 47,79% respectively.
- Less performant results for “authors” and “subjects” columns:
 - **High Cardinality:** a vast amount of unique labels to predict.
 - **Lack of Standardization:** the level of detail is different, particularly for “subjects” column (detailed vs. generic).

POSSIBLE FUTURE DIRECTIONS

- Exploring alternative training strategies and loss functions
- Trying to create a standardized method for describing an artwork.

THANK YOU